

Cognitive Rhythm as an Oversight Interface: Stop Rubber-Stamping through State-Aware Interaction

Ning Coeva

Independent Researcher

Accepted contribution to the AI CHAOS! Workshop, CHI 2026

Archival note. This author-prepared PDF reproduces the substantive text submitted through the workshop's Google Form. It was created after acceptance to provide a stable public copy; administrative form fields are omitted.

1. How would you define “oversight”?

Oversight is not a permission layer—it is a state-dependent capability. Meaningful oversight exists only when a human can, at the moment of a consequential decision: (i) form an understanding of what the system proposes and why, (ii) consider alternatives or challenge the output against evidence, and (iii) intervene to modify, delay, escalate, or reject. A click made without cognitive readiness is not oversight but compliance theater. Oversight therefore requires both an interface that supports sensemaking and a human in an oversight-ready state.

2. What activities should be considered oversight?

Activities qualify as oversight when they preserve a triad: visibility → sensemaking → agency. The human must access a meaningful representation of the system's proposed action, evaluate it against evidence or alternatives, and change the outcome if needed. Examples: reviewing an AI diagnosis with supporting evidence before treatment; inspecting an autonomous plan with override capability; auditing a hiring recommendation with enough context to contest it. The key criterion is not observation, but the ability to alter the system's trajectory.

3. What activities should not be considered oversight?

Activities should not count as oversight when the human cannot meaningfully alter outcomes or realistically evaluate the decision in time. This includes: rubber-stamping under time pressure, post-hoc notifications after actions are taken, interfaces too complex or fast to process, and approvals in depleted states (fatigue, overload). I call these Drive-state approvals: approvals in task-execution mode where the default becomes acceptance. They may satisfy logs, but create an illusion of control. Rubber-stamping is not oversight; it is oversight debt accumulating.

4. Example of effective oversight

A physician receives an AI-generated differential diagnosis showing alternatives with key supporting evidence. The interface supports quick sensemaking—evidence summaries and lightweight counterfactual checks—so she can test assumptions rather than accept a single output. Importantly, the decision is not forced under immediate time pressure, allowing an oversight-ready state. She notices an underweighted family-history factor, adjusts inputs, and the recommendation shifts. She proceeds with a modified plan. This is effective oversight: actionable visibility, time to reason, and real agency to change the outcome.

5. Example of ineffective oversight

A radiologist reviews AI-flagged X-rays at the end of a long shift. The interface requires confirm/reject clicks to proceed. Facing a backlog and fatigue, she slips into Drive mode—task-completion focus rather than critical evaluation—clicking through rapidly. The audit trail shows “100% human review,” but meaningful oversight occurred sporadically at best. This is the oversight illusion: formal compliance masking substantive failure. The system demanded oversight at the moment the human was least able to provide it.

6. What should everyone agree upon?

Everyone should agree that oversight effectiveness depends on human cognitive readiness, not just interface affordances. It is not enough that a human “could have intervened”—we must ask whether they were realistically capable of meaningful intervention in that moment. This shifts design from “did we show information?” to “did we request oversight when the human could actually provide it?” Oversight is a capability, not a checkbox.

7. What is the most pressing challenge?

The most pressing challenge is preventing oversight from degrading into performative compliance through oversight fatigue. As AI scales, humans face increasing oversight requests; without careful design, this accumulates oversight debt—approvals without genuine evaluation—until rubber-stamping becomes the norm. Systems must treat oversight as a scarce resource: budget and prioritize oversight moments, consolidate low-risk checks, and gate high-stakes commitments to oversight-ready states rather than task-execution pressure. Cognitive rhythm must be treated as part of the oversight interface.

8. Spicy take

Drive-state approvals should be treated as a safety smell, not evidence of compliance. When approvals occur under time pressure, high load, or task-completion focus, they should raise red flags rather than satisfy checkboxes. Systems should refuse irreversible confirmations when contextual